

## EDA Processing

```
In [21]: import pandas as pd
import numpy as np
import seaborn as sns
```

```
In [22]: data = pd.read_csv(r'C:\Users\HOME\Desktop\DSA\Lab6\data_clean.csv')
#data.describe
data
```

```
Out[22]:
```

	Unnamed: 0	Ozone	Solar.R	Wind	Temp C	Month	Day	Year	Temp	Weather
0	1	41.0	190.0	7.4	67	5	1	2010	67	S
1	2	36.0	118.0	8.0	72	5	2	2010	72	C
2	3	12.0	149.0	12.6	74	5	3	2010	74	PS
3	4	18.0	313.0	11.5	62	5	4	2010	62	S
4	5	NaN	NaN	14.3	56	5	5	2010	56	S
...	...	...	...	...	...	...	...	...	...	...
153	154	41.0	190.0	7.4	67	5	1	2010	67	C
154	155	30.0	193.0	6.9	70	9	26	2010	70	PS
155	156	NaN	145.0	13.2	77	9	27	2010	77	S
156	157	14.0	191.0	14.3	75	9	28	2010	75	S
157	158	18.0	131.0	8.0	76	9	29	2010	76	C

158 rows × 10 columns

```
In [23]: data.tail(10)
```

```
Out[23]:
```

	Unnamed: 0	Ozone	Solar.R	Wind	Temp C	Month	Day	Year	Temp	Weather
148	149	30.0	193.0	6.9	70	9	26	2010	70	C
149	150	NaN	145.0	13.2	77	9	27	2010	77	PS
150	151	14.0	191.0	14.3	75	9	28	2010	75	S
151	152	18.0	131.0	8.0	76	9	29	2010	76	PS
152	153	20.0	223.0	11.5	68	9	30	2010	68	S
153	154	41.0	190.0	7.4	67	5	1	2010	67	C
154	155	30.0	193.0	6.9	70	9	26	2010	70	PS
155	156	NaN	145.0	13.2	77	9	27	2010	77	S
156	157	14.0	191.0	14.3	75	9	28	2010	75	S
157	158	18.0	131.0	8.0	76	9	29	2010	76	C

```
In [24]: type(data)
data.shape
```

```
Out[24]: (158, 10)
```

```
In [25]: data.dtypes
```

```
Out[25]: Unnamed: 0      int64
Ozone      float64
Solar.R    float64
Wind       float64
Temp C     object
Month      object
Day        int64
Year       int64
Temp       int64
Weather    object
dtype: object
```

## Data Type Conversion

```
In [26]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 158 entries, 0 to 157
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Unnamed: 0   158 non-null   int64
1   Ozone        120 non-null   float64
2   Solar.R     151 non-null   float64
3   Wind         158 non-null   float64
4   Temp C       158 non-null   object
5   Month        158 non-null   object
6   Day          158 non-null   int64
7   Year         158 non-null   int64
8   Temp         158 non-null   int64
9   Weather      155 non-null   object
dtypes: float64(3), int64(4), object(3)
memory usage: 12.5+ KB
```

In [27]: *#checking missing values*

```
data2 = data.iloc[:,1:]  
data2
```

Out[27]:

	Ozone	Solar.R	Wind	Temp C	Month	Day	Year	Temp	Weather
0	41.0	190.0	7.4	67	5	1	2010	67	S
1	36.0	118.0	8.0	72	5	2	2010	72	C
2	12.0	149.0	12.6	74	5	3	2010	74	PS
3	18.0	313.0	11.5	62	5	4	2010	62	S
4	NaN	NaN	14.3	56	5	5	2010	56	S
...	...	...	...	...	...	...	...	...	...
153	41.0	190.0	7.4	67	5	1	2010	67	C
154	30.0	193.0	6.9	70	9	26	2010	70	PS
155	NaN	145.0	13.2	77	9	27	2010	77	S
156	14.0	191.0	14.3	75	9	28	2010	75	S
157	18.0	131.0	8.0	76	9	29	2010	76	C

158 rows × 9 columns

In [28]: *# .copy is used so any changes made in this wont be reflected in original data*

```
data = data2.copy()  
data
```

Out[28]:

	Ozone	Solar.R	Wind	Temp C	Month	Day	Year	Temp	Weather
0	41.0	190.0	7.4	67	5	1	2010	67	S
1	36.0	118.0	8.0	72	5	2	2010	72	C
2	12.0	149.0	12.6	74	5	3	2010	74	PS
3	18.0	313.0	11.5	62	5	4	2010	62	S
4	NaN	NaN	14.3	56	5	5	2010	56	S
...	...	...	...	...	...	...	...	...	...
153	41.0	190.0	7.4	67	5	1	2010	67	C
154	30.0	193.0	6.9	70	9	26	2010	70	PS
155	NaN	145.0	13.2	77	9	27	2010	77	S
156	14.0	191.0	14.3	75	9	28	2010	75	S
157	18.0	131.0	8.0	76	9	29	2010	76	C

158 rows × 9 columns

```
In [30]: data['Month'] = pd.to_numeric(data['Month'],errors='coerce')
data['Temp C'] = pd.to_numeric(data['Temp C'],errors='coerce')
data['Weather'] = data['Weather'].astype('category')
```

```
In [40]: #dropping temp column
data = data.drop('Temp',axis=1,inplace=True)
data
```

```
Out[40]:
```

	Ozone	Solar.R	Wind	Temp C	Month	Day	Year	Weather
0	41.0	190.0	7.4	67.0	5.0	1	2010	S
1	36.0	118.0	8.0	72.0	5.0	2	2010	C
2	12.0	149.0	12.6	74.0	5.0	3	2010	PS
3	18.0	313.0	11.5	62.0	5.0	4	2010	S
4	NaN	NaN	14.3	56.0	5.0	5	2010	S
...	...	...	...	...	...	...	...	...
153	41.0	190.0	7.4	67.0	5.0	1	2010	C
154	30.0	193.0	6.9	70.0	9.0	26	2010	PS
155	NaN	145.0	13.2	77.0	9.0	27	2010	S
156	14.0	191.0	14.3	75.0	9.0	28	2010	S
157	18.0	131.0	8.0	76.0	9.0	29	2010	C

158 rows × 8 columns

```
In [47]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 158 entries, 0 to 157
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Ozone       120 non-null    float64
1   Solar.R     151 non-null    float64
2   Wind        158 non-null    float64
3   Temp C      157 non-null    float64
4   Month       157 non-null    float64
5   Day         158 non-null    int64
6   Year        158 non-null    int64
7   Weather     155 non-null    category
dtypes: category(1), float64(5), int64(2)
memory usage: 9.1 KB
```

## EDA Processing Starts

### Finding Duplicates

```
In [48]: # Identifying duplicate values from data
```

```
data[data.duplicated()].shape
```

```
#duplicate
```

```
Out[48]: (1, 8)
```

```
In [50]: #printing duplicated row
```

```
data[data.duplicated()]
```

```
Out[50]:
```

	Ozone	Solar.R	Wind	Temp C	Month	Day	Year	Weather
156	14.0	191.0	14.3	75.0	9.0	28	2010	S

```
In [55]: #dropping duplicated rows
```

```
data_cleaned1 = data.drop_duplicates()
```

```
data_cleaned1.shape
```

```
Out[55]: (157, 8)
```

## Drop columns

```
In [57]: # dropping duplicated temp column no need to do since already done hence copyi
#data_cleaned2 = data_cleaned1.drop('Temp C',axis = 1)
data_cleaned2 = data_cleaned1.copy()
data_cleaned2
```

```
Out[57]:
```

	Ozone	Solar.R	Wind	Temp C	Month	Day	Year	Weather
0	41.0	190.0	7.4	67.0	5.0	1	2010	S
1	36.0	118.0	8.0	72.0	5.0	2	2010	C
2	12.0	149.0	12.6	74.0	5.0	3	2010	PS
3	18.0	313.0	11.5	62.0	5.0	4	2010	S
4	NaN	NaN	14.3	56.0	5.0	5	2010	S
...	...	...	...	...	...	...	...	...
152	20.0	223.0	11.5	68.0	9.0	30	2010	S
153	41.0	190.0	7.4	67.0	5.0	1	2010	C
154	30.0	193.0	6.9	70.0	9.0	26	2010	PS
155	NaN	145.0	13.2	77.0	9.0	27	2010	S
157	18.0	131.0	8.0	76.0	9.0	29	2010	C

157 rows × 8 columns

## Rename column

```
In [70]: data_cleaned3 = data_cleaned2.rename({'Solar.R': 'Solar'}, axis=1)
data_cleaned3
```

```
Out[70]:
```


	Ozone	Solar	Wind	Temp C	Month	Day	Year	Weather
0	41.0	190.0	7.4	67.0	5.0	1	2010	S
1	36.0	118.0	8.0	72.0	5.0	2	2010	C
2	12.0	149.0	12.6	74.0	5.0	3	2010	PS
3	18.0	313.0	11.5	62.0	5.0	4	2010	S
4	NaN	NaN	14.3	56.0	5.0	5	2010	S
...	...	...	...	...	...	...	...	...
152	20.0	223.0	11.5	68.0	9.0	30	2010	S
153	41.0	190.0	7.4	67.0	5.0	1	2010	C
154	30.0	193.0	6.9	70.0	9.0	26	2010	PS
155	NaN	145.0	13.2	77.0	9.0	27	2010	S
157	18.0	131.0	8.0	76.0	9.0	29	2010	C

157 rows × 8 columns

## Outlier Detection


```
In [67]: #histogram of ozone  
data_cleaned3['Ozone'].hist()
```

Out[67]: <Axes: >



```
In [65]: #boxplot
data_cleaned3.boxplot(column=['Ozone'])
```

Out[65]: <Axes: >



```
In [71]: #Descriptive stat
data_cleaned3['Ozone'].describe()
```

Out[71]:

count	119.000000
mean	41.815126
std	32.659249
min	1.000000
25%	18.000000
50%	31.000000
75%	62.000000
max	168.000000

Name: Ozone, dtype: float64



```
In [72]: data_cleaned3
```


```
Out[72]:
```

	Ozone	Solar	Wind	Temp C	Month	Day	Year	Weather
0	41.0	190.0	7.4	67.0	5.0	1	2010	S
1	36.0	118.0	8.0	72.0	5.0	2	2010	C
2	12.0	149.0	12.6	74.0	5.0	3	2010	PS
3	18.0	313.0	11.5	62.0	5.0	4	2010	S
4	NaN	NaN	14.3	56.0	5.0	5	2010	S
...	...	...	...	...	...	...	...	...
152	20.0	223.0	11.5	68.0	9.0	30	2010	S
153	41.0	190.0	7.4	67.0	5.0	1	2010	C
154	30.0	193.0	6.9	70.0	9.0	26	2010	PS
155	NaN	145.0	13.2	77.0	9.0	27	2010	S
157	18.0	131.0	8.0	76.0	9.0	29	2010	C

157 rows × 8 columns

```
In [73]: #bar plot
data['Weather'].value_counts().plot.bar()
```

```
Out[73]: <Axes: xlabel='Weather'>
```




## Missing Values and Imputation

```
In [82]: import seaborn as sns

cols = data_cleaned3.columns
colours = ['#000099', '#ffff00']
sns.heatmap(data_cleaned3[cols].isnull(), cmap=sns.color_palette(colours))
```

Out[82]: <Axes: >



```
In [83]: data_cleaned3[data_cleaned3.isnull().any(axis=1)].head()
```

Out[83]:

	Ozone	Solar	Wind	Temp C	Month	Day	Year	Weather
4	NaN	NaN	14.3	56.0	5.0	5	2010	S
5	28.0	NaN	14.9	66.0	5.0	6	2010	C
9	NaN	194.0	8.6	69.0	5.0	10	2010	S
10	7.0	NaN	6.9	NaN	5.0	11	2010	C
23	32.0	92.0	12.0	61.0	NaN	24	2010	C

```
In [84]: data_cleaned3.isnull().sum()
```

```
Out[84]: Ozone      38
          Solar      7
          Wind       0
          Temp C     1
          Month      1
          Day        0
          Year       0
          Weather    3
          dtype: int64
```

```
In [85]: #Mean Imputation
```

```
mean = data_cleaned3['Ozone'].mean()
mean
```

```
Out[85]: 41.81512605042017
```

```
In [87]: data_cleaned3['Ozone'] = data_cleaned3['Ozone'].fillna(mean)
data_cleaned3
```

```
Out[87]:
```

	Ozone	Solar	Wind	Temp C	Month	Day	Year	Weather
0	41.000000	190.0	7.4	67.0	5.0	1	2010	S
1	36.000000	118.0	8.0	72.0	5.0	2	2010	C
2	12.000000	149.0	12.6	74.0	5.0	3	2010	PS
3	18.000000	313.0	11.5	62.0	5.0	4	2010	S
4	41.815126	NaN	14.3	56.0	5.0	5	2010	S
...	...	...	...	...	...	...	...	...
152	20.000000	223.0	11.5	68.0	9.0	30	2010	S
153	41.000000	190.0	7.4	67.0	5.0	1	2010	C
154	30.000000	193.0	6.9	70.0	9.0	26	2010	PS
155	41.815126	145.0	13.2	77.0	9.0	27	2010	S
157	18.000000	131.0	8.0	76.0	9.0	29	2010	C

157 rows × 8 columns

```
In [88]: #Mean Imputation of solar
```

```
mean = data_cleaned3['Solar'].mean()
data_cleaned3['Solar'] = data_cleaned3['Solar'].fillna(mean)
mean
```

```
Out[88]: 185.36666666666667
```

In [90]: *#Mean Imputation of Month*

```
mean = data_cleaned3['Month'].mean()
data_cleaned3['Month'] = data_cleaned3['Month'].fillna(mean)
mean
```

Out[90]: 7.032051282051282

In [92]: *#Mean Imputation of Temp C*

```
mean = data_cleaned3['Temp C'].mean()
data_cleaned3['Temp C'] = data_cleaned3['Temp C'].fillna(mean)
mean
```

Out[92]: 77.76923076923077

In [94]: *# Imputation of Weather*

```
obj_columns = data_cleaned3[['Weather']]
obj_columns.isnull().sum()
```

Out[94]: Weather 3  
dtype: int64

In [95]: obj\_columns

Out[95]:

	Weather
0	S
1	C
2	PS
3	S
4	S
...	...
152	S
153	C
154	PS
155	S
157	C

157 rows × 1 columns

In [96]: *#missing value imputation for categorical value*

```
obj_columns = obj_columns.fillna(obj_columns.mode().iloc[0])
```

```
In [97]: obj_columns.mode()
```

```
Out[97]:
```

	Weather
0	S

```
In [99]: obj_columns.isnull().sum()
```

```
Out[99]: Weather    0
dtype: int64
```

```
In [102]: data_cleaned3.shape
```

```
Out[102]: (157, 8)
```

```
In [103]: obj_columns.shape
```

```
Out[103]: (157, 1)
```

```
In [105]: #join data set with imputed value
```

```
data_cleaned4 = pd.concat([data_cleaned3,obj_columns],axis=1)
```

```
In [108]: data_cleaned4
```

```
Out[108]:
```

	Ozone	Solar	Wind	Temp C	Month	Day	Year	Weather	Weather
0	41.000000	190.000000	7.4	67.0	5.0	1	2010	S	S
1	36.000000	118.000000	8.0	72.0	5.0	2	2010	C	C
2	12.000000	149.000000	12.6	74.0	5.0	3	2010	PS	PS
3	18.000000	313.000000	11.5	62.0	5.0	4	2010	S	S
4	41.815126	185.366667	14.3	56.0	5.0	5	2010	S	S
...	...	...	...	...	...	...	...	...	...
152	20.000000	223.000000	11.5	68.0	9.0	30	2010	S	S
153	41.000000	190.000000	7.4	67.0	5.0	1	2010	C	C
154	30.000000	193.000000	6.9	70.0	9.0	26	2010	PS	PS
155	41.815126	145.000000	13.2	77.0	9.0	27	2010	S	S
157	18.000000	131.000000	8.0	76.0	9.0	29	2010	C	C

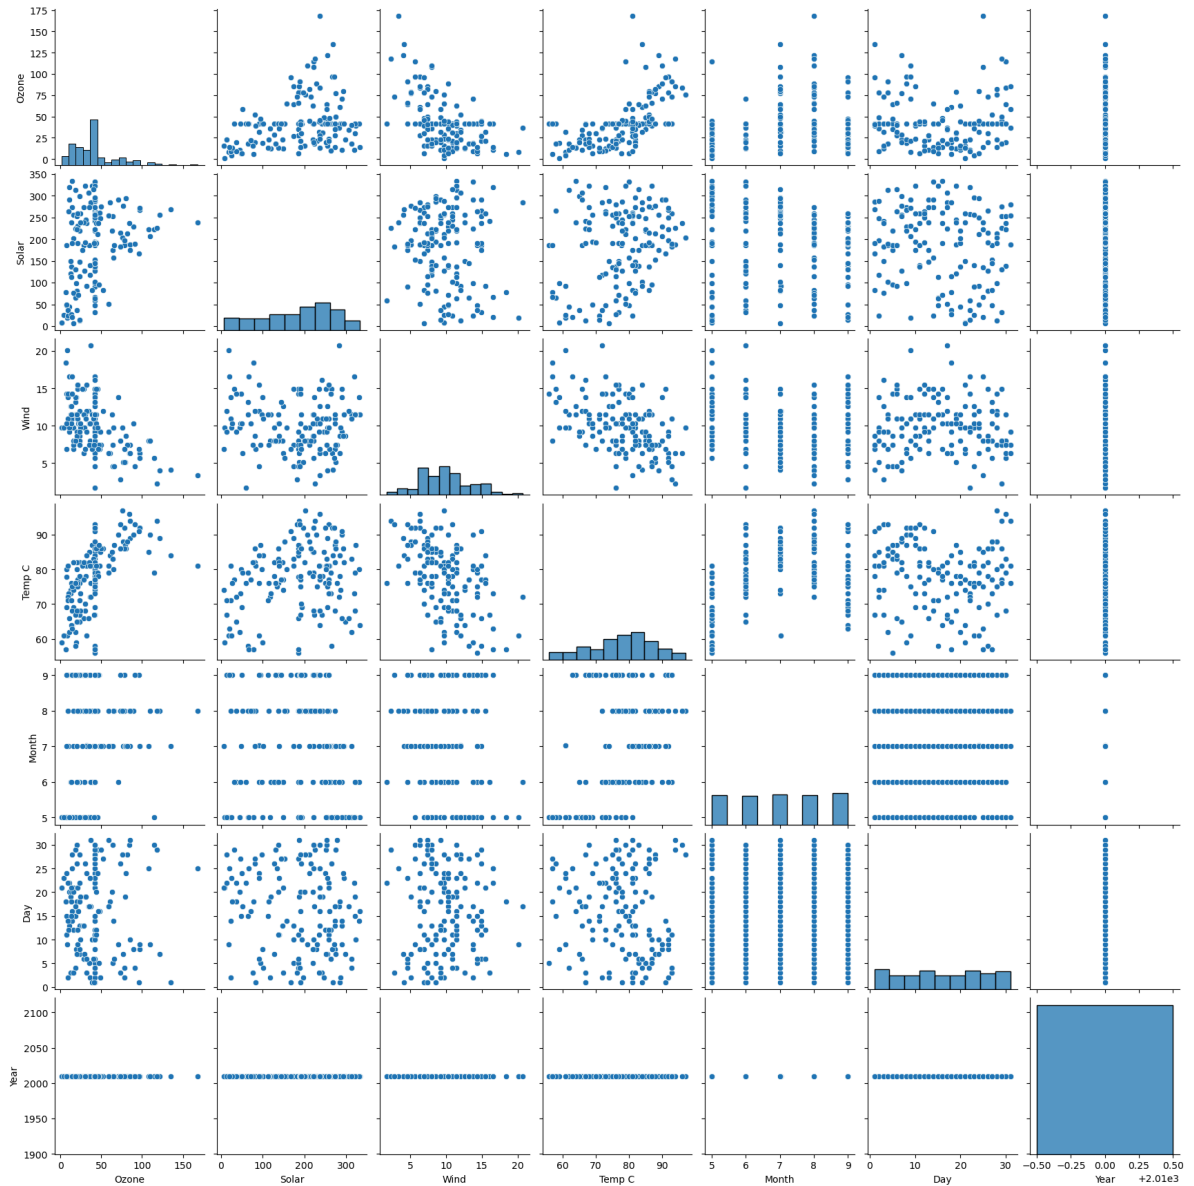
157 rows × 9 columns

## Scatter Plot and correlation analysis

```
In [109]: sns.pairplot(data_cleaned3)
```

```
c:\Users\HOME\anaconda3\Lib\site-packages\seaborn\axisgrid.py:123: UserWarning: The figure layout has changed to tight
  self._figure.tight_layout(*args, **kwargs)
```

```
Out[109]: <seaborn.axisgrid.PairGrid at 0x22c87856d90>
```



```
In [111]: # Select only numeric columns
numeric_data = data_cleaned3.select_dtypes(include=np.number)

#compute correlation matrix
correlation_matrix = numeric_data.corr()

#print correlation matrix
correlation_matrix
```

```
Out[111]:
```

	Ozone	Solar	Wind	Temp C	Month	Day	Year
Ozone	1.000000	0.304559	-0.520004	0.603660	0.132809	-0.021916	NaN
Solar	0.304559	1.000000	-0.055874	0.260810	-0.090564	-0.151007	NaN
Wind	-0.520004	-0.055874	1.000000	-0.443676	-0.166029	0.029900	NaN
Temp C	0.603660	0.260810	-0.443676	1.000000	0.390957	-0.124262	NaN
Month	0.132809	-0.090564	-0.166029	0.390957	1.000000	0.049924	NaN
Day	-0.021916	-0.151007	0.029900	-0.124262	0.049924	1.000000	NaN
Year	NaN	NaN	NaN	NaN	NaN	NaN	NaN

## Transformations

### Dummy Variable

```
In [115]: #creating dummy for weather

data_cleaned4 = pd.get_dummies(data,columns=['Weather'])
data_cleaned4
```

```
Out[115]:
```

	Ozone	Solar.R	Wind	Temp C	Month	Day	Year	Weather_C	Weather_PS	Weather_S
0	41.0	190.0	7.4	67.0	5.0	1	2010	False	False	True
1	36.0	118.0	8.0	72.0	5.0	2	2010	True	False	False
2	12.0	149.0	12.6	74.0	5.0	3	2010	False	True	False
3	18.0	313.0	11.5	62.0	5.0	4	2010	False	False	True
4	NaN	NaN	14.3	56.0	5.0	5	2010	False	False	True
...	...	...	...	...	...	...	...	...	...	...
153	41.0	190.0	7.4	67.0	5.0	1	2010	True	False	False
154	30.0	193.0	6.9	70.0	9.0	26	2010	False	True	False
155	NaN	145.0	13.2	77.0	9.0	27	2010	False	False	True
156	14.0	191.0	14.3	75.0	9.0	28	2010	False	False	True
157	18.0	131.0	8.0	76.0	9.0	29	2010	True	False	False

158 rows × 10 columns

```
In [116]: data_cleaned4 = data_cleaned4.dropna()
```

## Normalization of data

```
In [117]: from numpy import set_printoptions
from sklearn.preprocessing import MinMaxScaler
```

```
In [118]: data_cleaned4.values
```

```
Out[118]: array([[41.0, 190.0, 7.4, ..., False, False, True],
                [36.0, 118.0, 8.0, ..., True, False, False],
                [12.0, 149.0, 12.6, ..., False, True, False],
                ...,
                [30.0, 193.0, 6.9, ..., False, True, False],
                [14.0, 191.0, 14.3, ..., False, False, True],
                [18.0, 131.0, 8.0, ..., True, False, False]], dtype=object)
```

```
In [120]: array = data_cleaned3.values
scaler = MinMaxScaler(feature_range=(0,1))
rescaledX = scaler.fit_transform(array[:,0:5])
```

```
#transformed data
set_printoptions(precision=2)
print(rescaledX[0:5,:])
```

```
[[0.24 0.56 0.3  0.27 0. ]
 [0.21 0.34 0.33 0.39 0. ]
 [0.07 0.43 0.57 0.44 0. ]
 [0.1  0.94 0.52 0.15 0. ]
 [0.24 0.55 0.66 0.   0. ]]
```

```
In [121]: #standardize data (0mean,1sd)
from sklearn.preprocessing import StandardScaler
```

```
In [122]: array = data_cleaned4.values
scaler = StandardScaler().fit(array)
rescaledX = scaler.transform(array)
```

```
#summarize
set_printoptions(precision = 2)
print(rescaledX[0:5,:])
```

```
[[ -0.02  0.05 -0.71 -1.15 -1.53 -1.7   0.   -0.64 -0.68  1.28]
 [ -0.17 -0.75 -0.54 -0.62 -1.53 -1.59  0.    1.57 -0.68 -0.78]
 [ -0.9  -0.41  0.77 -0.4  -1.53 -1.48  0.   -0.64  1.47 -0.78]
 [ -0.72  1.43  0.45 -1.69 -1.53 -1.36  0.   -0.64 -0.68  1.28]
 [ -0.57  1.27 -0.37 -1.37 -1.53 -1.02  0.   -0.64  1.47 -0.78]]
```



## Speed up the EDA process

```
In [123]: #install dtale  
  
import dtale  
dtale.show(data)
```

Out[123]: