

Lab 7 Multiple Linear Regression

```
In [3]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.graphics.regressionplots import influence_plot
import statsmodels.formula.api as smf
import numpy as np
```

```
In [4]: # read data
cars = pd.read_csv('Cars.csv')
cars.head()
```

```
Out[4]:
```

	HP	MPG	VOL	SP	WT
0	49	53.700681	89	104.185353	28.762059
1	55	50.013401	92	105.461264	30.466833
2	55	50.013401	92	105.461264	30.193597
3	70	45.696322	92	113.461264	30.632114
4	53	50.504232	92	104.461264	29.889149

```
In [5]: cars.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 81 entries, 0 to 80
Data columns (total 5 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   HP      81 non-null     int64  
 1   MPG     81 non-null     float64
 2   VOL     81 non-null     int64  
 3   SP      81 non-null     float64
 4   WT      81 non-null     float64
dtypes: float64(3), int64(2)
memory usage: 3.3 KB
```

```
In [6]: #checking null
cars.isna().sum()
```

```
Out[6]: HP      0
MPG      0
VOL      0
SP       0
WT       0
dtype: int64
```

```
In [7]: cars.corr()
```

```
Out[7]:
```

	HP	MPG	VOL	SP	WT
HP	1.000000	-0.725038	0.077459	0.973848	0.076513
MPG	-0.725038	1.000000	-0.529057	-0.687125	-0.526759
VOL	0.077459	-0.529057	1.000000	0.102170	0.999203
SP	0.973848	-0.687125	0.102170	1.000000	0.102439
WT	0.076513	-0.526759	0.999203	0.102439	1.000000


Scatterplot between variables along histograms

```
In [8]: # format plot background and scatter plot for all variable
```

```
sns.set_style(style='darkgrid')  
sns.pairplot(cars)
```

```
c:\Users\HOME\anaconda3\Lib\site-packages\seaborn\axisgrid.py:123: UserWarning:  
The figure layout has changed to tight  
self._figure.tight_layout(*args, **kwargs)
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x24b7ba0b4d0>
```



Preparing model

```
In [12]: import statsmodels.formula.api as smf

# Assuming 'cars' is your DataFrame containing the data
model = smf.ols('MPG ~ WT + VOL + SP + HP', data=cars).fit()
```

```
In [13]: model.params
```

```
Out[13]: Intercept    30.677336
         WT           0.400574
         VOL          -0.336051
         SP           0.395627
         HP          -0.205444
         dtype: float64
```

```
In [14]: # t and p values
print(model.tvalues, '\n', model.pvalues)
```

```
Intercept    2.058841
WT            0.236541
VOL          -0.590970
SP            2.499880
HP           -5.238735
dtype: float64
Intercept    0.042936
WT            0.813649
VOL            0.556294
SP            0.014579
HP            0.000001
dtype: float64
```

```
In [15]: #R squared values
(model.rsquared, model.rsquared_adj)
```

```
Out[15]: (0.7705372737359844, 0.7584602881431415)
```

Simple Linear Regression

```
In [18]: import statsmodels.formula.api as smf

# Fit the linear regression model
ml_v = smf.ols('MPG ~ VOL', data=cars).fit()

# Print t-values and p-values
print(ml_v.tvalues, '\n', ml_v.pvalues)
```

```
Intercept    14.106056
VOL          -5.541400
dtype: float64
Intercept    2.753815e-23
VOL          3.822819e-07
dtype: float64
```

```
In [21]: ml_w = smf.ols('MPG ~ WT', data=cars).fit()
print(ml_w.tvalues, '\n', ml_w.pvalues)
```

```
Intercept    14.248923
WT           -5.508067
dtype: float64
Intercept    1.550788e-23
WT           4.383467e-07
dtype: float64
```

```
In [22]: ml_wv = smf.ols('MPG~WT+VOL', data=cars).fit()
print(ml_wv.tvalues, '\n', ml_wv.pvalues)
```

```
Intercept    12.545736
WT            0.489876
VOL          -0.709604
dtype: float64
Intercept    2.141975e-20
WT           6.255966e-01
VOL           4.800657e-01
dtype: float64
```

Calculating VIF

```
In [28]: rsq_hp = smf.ols('HP~WT+VOL+SP',data=cars).fit().rsquared
vif_hp = 1/(1-rsq_hp) # 16.33

rsq_wt = smf.ols('WT~HP+VOL+SP',data=cars).fit().rsquared
vif_wt = 1/(1-rsq_wt) # 564.98

rsq_vol = smf.ols('VOL~HP+WT+SP',data=cars).fit().rsquared
vif_vol = 1/(1-rsq_vol) # 16.33

rsq_sp = smf.ols('SP~WT+VOL+HP',data=cars).fit().rsquared
vif_sp = 1/(1-rsq_sp) # 16.33

#storing vif values in dataframe
d1 = {'Variables':['HP','WT','VOL','SP'],'VIF':[vif_hp,vif_wt,vif_vol,vif_sp]}
d1_frame = pd.DataFrame(d1)
d1_frame
```


Out[28]:

	Variables	VIF
0	HP	19.926589
1	WT	639.533818
2	VOL	638.806084
3	SP	20.007639

Residual Analysis

Test for normality of residuals (Q-Q plot)

```
In [30]: import statsmodels.api as sm
qqplot = sm.qqplot(model.resid, line='q') # line = 45 to draw diagonal line
plt.title('Normal Q-Q plot of residuals')
plt.show()
```



```
In [31]: list(np.where(model.resid>10))
```

```
Out[31]: [array([ 0, 76], dtype=int64)]
```

Residual Plot for Homoscedasticity

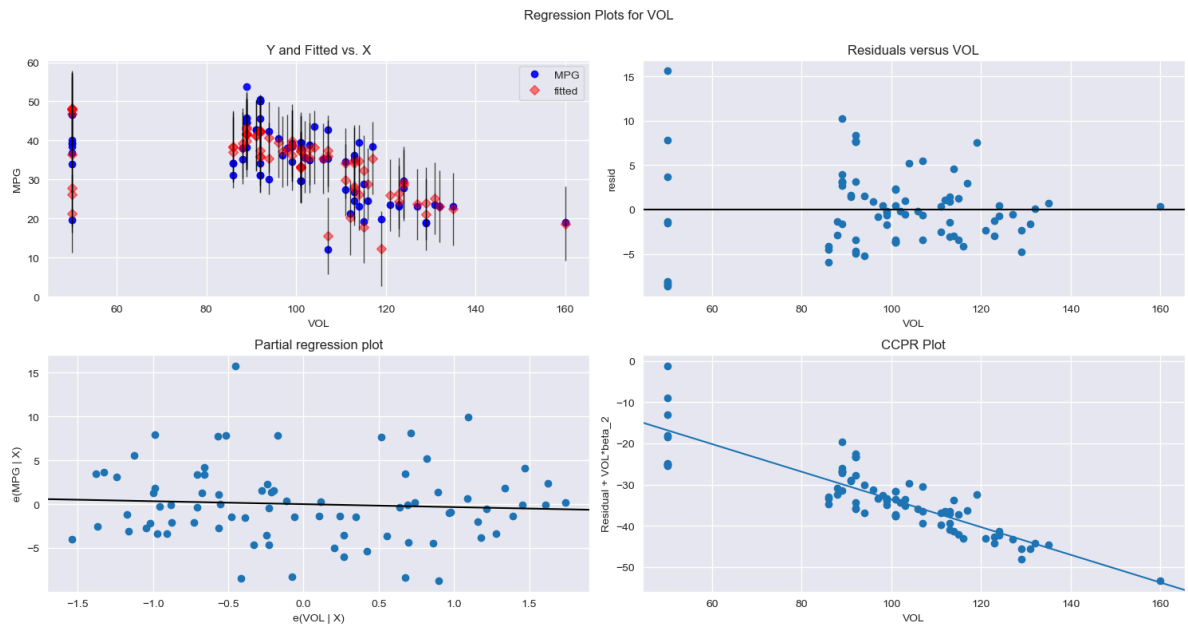
```
In [32]: def get_standardized_values(vals):
          return (vals-vals.mean())/vals.std()
```

```
In [33]: plt.scatter(get_standardized_values(model.fittedvalues),get_standardized_value
plt.title('Residual Plot')
plt.xlabel('Standardized Fitted Values')
plt.ylabel('Standardized Residual Values')
plt.show()
```

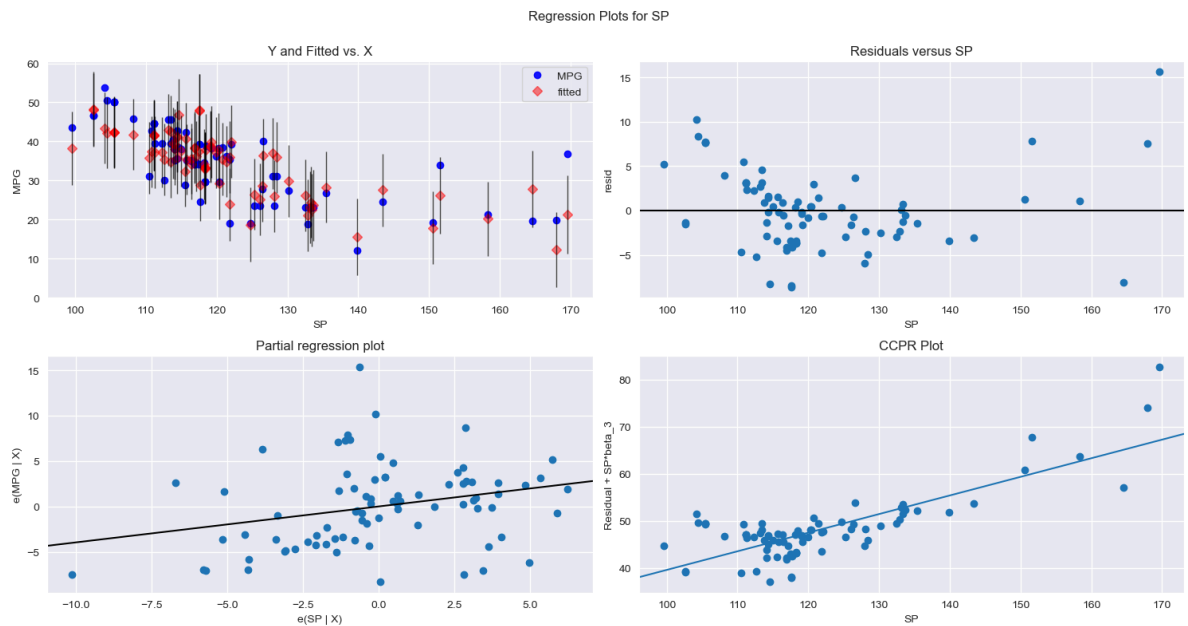


Residual VS Regressors

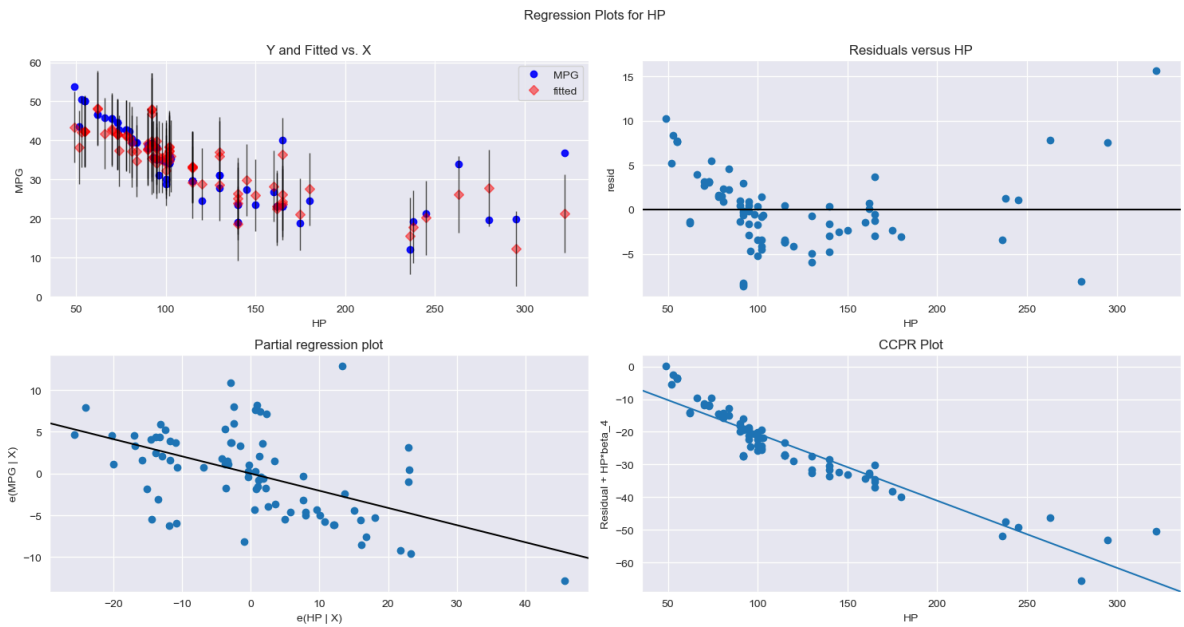
```
In [34]: fig = plt.figure(figsize=(15,8))
fig = sm.graphics.plot_regress_exog(model, 'VOL', fig=fig)
plt.show()
```



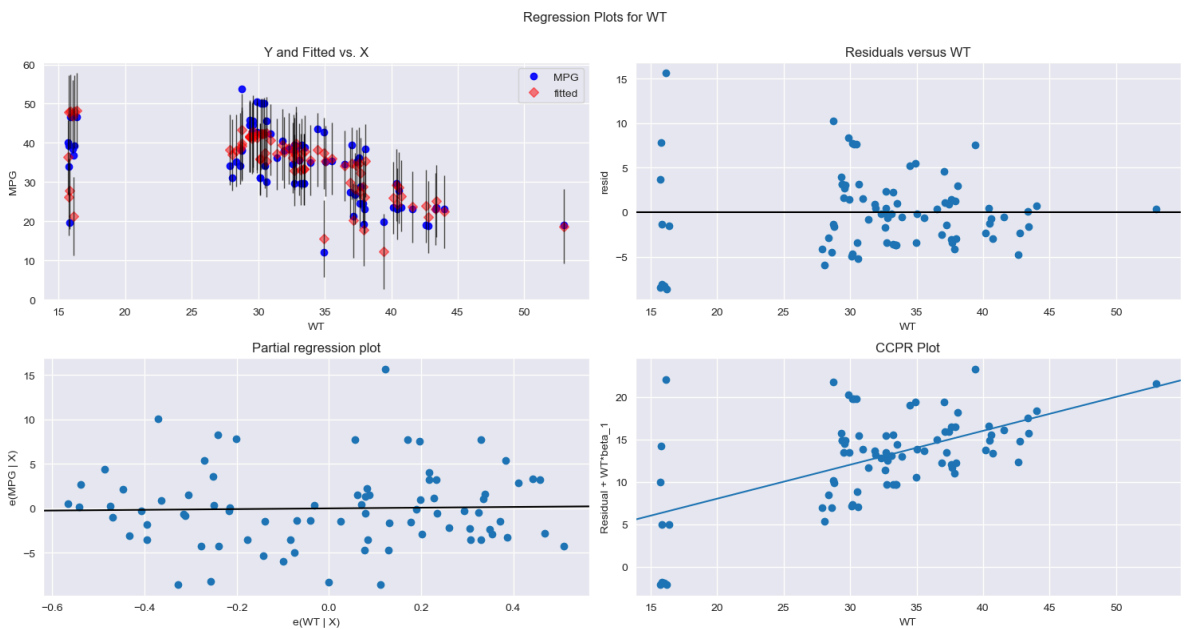
```
In [36]: fig = plt.figure(figsize=(15,8))
fig = sm.graphics.plot_regress_exog(model, 'SP', fig=fig)
plt.show()
```



```
In [37]: fig = plt.figure(figsize=(15,8))
fig = sm.graphics.plot_regress_exog(model, 'HP', fig=fig)
plt.show()
```



```
In [38]: fig = plt.figure(figsize=(15,8))
fig = sm.graphics.plot_regress_exog(model, 'WT', fig=fig)
plt.show()
```



Model Detection Diagnostics

Detecting Influences/Outliers


Cook's Distance

```
In [39]: model_influence = model.get_influence()  
(c,_) = model_influence.cooks_distance
```

```
In [42]: #plot influencers values using stem plot
```

```
fig = plt.subplots(figsize=(20,7))  
plt.stem(np.arange(len(cars)), np.round(c,3))  
plt.xlabel('Row Index')  
plt.ylabel('Cooks Distance')  
plt.show()
```

```
# distance greater than 1 then outlier
```




```
In [43]: # index and value of outlier  
(np.argmax(c), np.max(c))
```

```
Out[43]: (76, 1.0865193998180098)
```

High influence points

```
In [45]: from statsmodels.graphics.regressionplots import influence_plot
influence_plot(model)
plt.show()
```



```
In [46]: k = cars.shape[1]
n = cars.shape[0]
leverage_cutoff = 3*((k+1)/n)
```

```
In [47]: cars[cars.index.isin([70,76])]
```

```
Out[47]:
```

	HP	MPG	VOL	SP	WT
70	280	19.678507	50	164.598513	15.823060
76	322	36.900000	50	169.598513	16.132947

```
In [48]: # See differences in HP and other variable values
cars.head()
```

```
Out[48]:
```

	HP	MPG	VOL	SP	WT
0	49	53.700681	89	104.185353	28.762059
1	55	50.013401	92	105.461264	30.466833
2	55	50.013401	92	105.461264	30.193597
3	70	45.696322	92	113.461264	30.632114
4	53	50.504232	92	104.461264	29.889149

Improving Model

```
In [49]: # Load data
cars_new = pd.read_csv('Cars.csv')
```

```
In [50]: # Discard datapoints which are influences and reassign row number(reset_index(
car1 = cars_new.drop(cars.index[[70,76]],axis=0).reset_index()
```

```
In [51]: #drop original index
car1=car1.drop(['index'],axis=1)
```

```
In [52]: car1
```

```
Out[52]:
```

	HP	MPG	VOL	SP	WT
0	49	53.700681	89	104.185353	28.762059
1	55	50.013401	92	105.461264	30.466833
2	55	50.013401	92	105.461264	30.193597
3	70	45.696322	92	113.461264	30.632114
4	53	50.504232	92	104.461264	29.889149
...
74	175	18.762837	129	132.864163	42.778219
75	238	19.197888	115	150.576579	37.923113
76	263	34.000000	50	151.598513	15.769625
77	295	19.833733	119	167.944460	39.423099
78	236	12.101263	107	139.840817	34.948615

79 rows × 5 columns

Build Model

```
In [54]: #Exclude variable 'WI' and generate R-squared and AIC Values
import statsmodels.formula.api as smf

# Fit the linear regression model excluding 'WI'
final_ml_v = smf.ols('MPG ~ VOL + SP + HP', data=car1).fit()

# Get R-squared and AIC values
r_squared = final_ml_v.rsquared
aic = final_ml_v.aic


print("R-squared:", r_squared)
print("AIC:", aic)
```

```
R-squared: 0.8161692010376007
AIC: 446.11722639447726
```

Comparing Above rsquare and AIC

COOK's Distance

```
In [56]: model_influence_v = final_ml_v.get_influence()
(c_v,_) = model_influence_v.cooks_distance
plt.stem(np.arange(len(car1)), np.round(c_v, 3));
plt.xlabel('Row Index')
plt.ylabel('Cooks Distance');
```



```
In [57]: #index of data where c > 0.5
(np.argmax(c_v), np.max(c_v))
```

```
Out[57]: (76, 1.1629387469135182)
```

```
In [58]: #drop 76 and 77
car2 = car1.drop(car1.index[[76,77]],axis=0)
```

```
In [59]: car2
```

```
Out[59]:
```

	HP	MPG	VOL	SP	WT
0	49	53.700681	89	104.185353	28.762059
1	55	50.013401	92	105.461264	30.466833
2	55	50.013401	92	105.461264	30.193597
3	70	45.696322	92	113.461264	30.632114
4	53	50.504232	92	104.461264	29.889149
...
72	140	19.086341	160	124.715241	52.997752
73	140	19.086341	129	121.864163	42.618698
74	175	18.762837	129	132.864163	42.778219
75	238	19.197888	115	150.576579	37.923113
78	236	12.101263	107	139.840817	34.948615

77 rows × 5 columns

```
In [60]: #reset index
car3 = car2.reset_index()
```

```
In [61]: car4 = car3.drop(['index'],axis=1)
```

```
In [62]: car4
```

```
Out[62]:
```


	HP	MPG	VOL	SP	WT
0	49	53.700681	89	104.185353	28.762059
1	55	50.013401	92	105.461264	30.466833
2	55	50.013401	92	105.461264	30.193597
3	70	45.696322	92	113.461264	30.632114
4	53	50.504232	92	104.461264	29.889149
...
72	140	19.086341	160	124.715241	52.997752
73	140	19.086341	129	121.864163	42.618698
74	175	18.762837	129	132.864163	42.778219
75	238	19.197888	115	150.576579	37.923113
76	236	12.101263	107	139.840817	34.948615

77 rows × 5 columns


```
In [63]: #build model
final_ml_v=smf.ols('MPG~VOL+SP+HP',data=car4).fit()
```

```
In [64]: # Again check for influencers
model_influence_v =final_ml_v.get_influence()
(c_v,_) = model_influence_v.cooks_distance
```

```
In [65]: fig = plt.subplots(figsize=(20,7))
plt.stem(np.arange(len(car4)),np.round(c_v,3));
plt.xlabel('Row Index')
plt.ylabel('Cooks Distance');
```



```
In [66]: #index of data where c > 0.5
(np.argmax(c_v),np.max(c_v))
```

```
Out[66]: (65, 0.8774556986296786)
```

Since value less than 1 we can stop diagnostic process and finalize the model

```
In [67]: # check accuracy of the model
final_ml_v=smf.ols('MPG~VOL+SP+HP',data=car4).fit()
```

```
In [68]: (final_ml_v.rsquared,final_ml_v.aic)
```

```
Out[68]: (0.8669636111859063, 409.4153062719507)
```

Predicting for new data

```
In [69]: # new data for prediction
new_data = pd.DataFrame({'HP':40,'VOL':95,"SP":102,"WT":35},index =[1])
```

```
In [70]: final_ml_v.predict(new_data)
```

```
Out[70]: 1    46.035594
dtype: float64
```

```
In [71]: final_ml_v.predict(cars_new.iloc[0:5,])
```

```
Out[71]: 0    45.428872  
         1    43.992392  
         2    43.992392  
         3    43.508150  
         4    44.085858  
         dtype: float64
```

```
In [72]: pred_y=final_ml_v.predict(cars_new)
```

```
In [73]: pred_y
```

```
Out[73]: 0    45.428872  
         1    43.992392  
         2    43.992392  
         3    43.508150  
         4    44.085858  
         ...  
        76    7.165876  
        77   12.198598  
        78   14.908588  
        79    4.163958  
        80    9.161202  
         Length: 81, dtype: float64
```

```
In [ ]:
```